

Project 4
CS 4471

This programming project has two parts, listed below.

Create a Gitlab repository labeled Project4 inside the group I invited you to. Work on your project in your Gitlab repository. When finished, tag your latest commit with “Finished” and I will know to grade it. The timestamp of the “Finished” commit will be used to determine if you completed the project on time.

Notes:

1. For this homework you will implement the sender-initiated load balancing algorithm described in Sec. 8.2. Your output should exactly match the output given below.
- Use the given starter code in `balance/`.
 - In the example output below for test file `test3.dat`, the first table (not the descriptive text at the beginning) prints which processes are assigned to run when on each CPU. For example, CPU A will run process 0 from 0-7. It would like to run process 1 from 6-13, but that isn’t possible since they overlap. The final table shows how the processes were actually run. CPU B ended up running process 1.
 - The middle table describes states of CPUs and actions. If a CPU is underutilized, then “under” is printed. If it is overutilized, then “over” is printed. In the case of overutilization, the system attempts to find an underutilized CPU and transfer a process to it. For example, in the `test3.dat` example, process 3 is transferred from CPU B to CPU C in timestep 3.
 - You should make sure you understand the example output before starting to code.
 - You can assume that we will have at most 20 time steps.
 - It is recommended that you store CPU ids as integers. So CPU ‘A’ would be 0, ‘B’ would be 1, etc. You can compute this using `cpu_idx = cpu_id - ‘A’`.

```
> ./balancing test1.dat
CPU A creates process 0 at time 0 requiring 8 time units.
CPU A creates process 1 at time 6 requiring 8 time units.

  0  1  2  3  4  5  6  7  8  9 10 11 12 13 14 15 16 17 18 19
A: 0  0  0  0  0  0  0  0  0
A:           1  1  1  1  1  1  1

Time Proc State Action
6     A    over
7     A    over

  0  1  2  3  4  5  6  7  8  9 10 11 12 13 14 15 16 17 18 19
A: 0  0  0  0  0  0  0  0  0
A:           1  1  1  1  1  1  1  1
```

```
> ./balancing test2.dat
CPU A creates process 0 at time 0 requiring 8 time units.
CPU A creates process 1 at time 6 requiring 8 time units.
CPU B creates process 2 at time 0 requiring 4 time units.
CPU B creates process 3 at time 3 requiring 5 time units.
CPU B creates process 4 at time 10 requiring 6 time units.

  0  1  2  3  4  5  6  7  8  9 10 11 12 13 14 15 16 17 18 19
A: 0  0  0  0  0  0  0  0  0
A:           1  1  1  1  1  1  1
B: 2  2  2  2
B:           3  3  3  3  3
B:                     4  4  4  4  4  4

Time Proc State Action
3     B    over
6     A    over
7     A    over
9     B    under

  0  1  2  3  4  5  6  7  8  9 10 11 12 13 14 15 16 17 18 19
```

```

A: 0 0 0 0 0 0 0 0 0
A: 1 1 1 1 1 1 1 1
B: 2 2 2 2
B: 3 3 3 3 3
B: 4 4 4 4 4 4

> ./balancing test3.dat
CPU A creates process 0 at time 0 requiring 8 time units.
CPU A creates process 1 at time 6 requiring 8 time units.
CPU B creates process 2 at time 0 requiring 4 time units.
CPU B creates process 3 at time 3 requiring 5 time units.
CPU B creates process 4 at time 10 requiring 6 time units.
CPU C creates process 5 at time 0 requiring 1 time units.

0 1 2 3 4 5 6 7 8 9 10 11 12 13 14 15 16 17 18 19
A: 0 0 0 0 0 0 0 0 0
A: 1 1 1 1 1 1
B: 2 2 2 2
B: 3 3 3 3 3
B: 4 4 4 4 4 4
C: 5

Time Proc State Action
1 C under
2 C under
3 B over process 3 -> CPU C
4 B under
5 B under
6 A over process 1 -> CPU B
8 A under
8 C under
9 A under
9 C under
10 A under
10 B over process 4 -> CPU A
10 C under
11 C under
12 C under
13 C under
14 B under
14 C under
15 B under
15 C under

0 1 2 3 4 5 6 7 8 9 10 11 12 13 14 15 16 17 18 19
A: 0 0 0 0 0 0 0 0 0
A: 4 4 4 4 4 4
B: 2 2 2 2
B: 1 1 1 1 1 1 1
C: 5
C: 3 3 3 3 3

```

2. You will implement a reader for steganographic images. Start from the given source code in `stega/`.
- The skeleton code reads a .bmp file for you and places the file in `buffer`. See comments in `stega.cpp` for instructions on how to interpret the buffer.
 - The bitmap files are encoded as follows: each byte of image data (NOT header data, so be sure to skip the header as discussed in `stega.cpp`) has one bit of encoded data – the least significant bit. The first byte of data’s least significant bit is the first character’s least significant bit. The second byte of data contributes to the first character’s second-least significant bit. Here are two examples:

bytes of data	bits	ASCII	character
255 254 254 118 118 118 121 124	--> 01000001	--> 65	--> 'A'
200 198 202 199 153 211 223 188	--> 01111000	--> 120	--> 'x'

Note: The bytes of data contribute to the ASCII bits in *reverse* order. For example, in the first example above, the byte 255 contributes the last one in the bits; 124 contributes the very first zero in the bits.

- How do you know if the last bit in a byte is a zero or one? There are two ways: you can either use a bitwise AND, or you can take the number modulo 2.
- How can you construct an ASCII byte of data from a bunch of bits? Initialize the byte `b` to zero, then if, say, the third-to-last bit is a one, `b = b | (1 << 3)`, which left-shifts a one three places and ORs it with byte `b`.
- Test using the three included data files.
- **IMPORTANT!** To receive full credit you must decode a certain image found on the web. Go to www2.cose.isu.edu/~mcgrmic2 and look for an image that has two gears. Download the image and run it through your steganography decoder. Follow the secret instructions for the final 10% of the project.

Following is a sample output.

```
> ./stega hello.bmp
Hello world!
```